

Log-periodic oscillation in the all-heads coin game

Peter Pfaffelhuber

May 21, 2026

Abstract

We study the large- n behaviour of the all-heads coin game below the fair coin. Two phenomena are established. First, for every $p \in (0, \frac{1}{2})$ the value $b_{n,p}$ of the greedy strategy ALL is asymptotically *log-periodic*: it is asymptotic to a continuous, non-constant function of $\log n$ with period $\log(1/(1-p))$, and in particular does not converge. Second — and this is the main result — the *optimal* winning probability $w_{n,p}$ inherits this oscillation: for every $p \in (0, \frac{1}{2})$ the sequence $n \mapsto w_{n,p}$ is not eventually monotone and has infinitely many strict local maxima and minima. The proof reduces the statement to two one-sided assertions, treats the increasing side by an elementary Bellman estimate, and the decreasing side by an exact transfer to strategy ALL followed by a renormalisation argument that propagates the oscillation across geometric scales. This paper is the asymptotic companion to the exact and perturbative analysis of [1].

*Working draft. Steps still to be completed are flagged in the text by boxed **OPEN GAP** markers (Gaps A–D); Section 9 consolidates their status.*

1 Introduction

1.1 The game

Fix $p \in (0, 1)$ and write $q := 1 - p$. A player holds n coins, each of which lands heads independently with probability p . In each round all remaining coins are tossed; the player must set aside at least one coin showing heads, and loses if no coin shows heads. The game is won once every coin has been set aside. Writing $w_{n,p}$ for the optimal winning probability, the dynamic-programming principle gives the Bellman equation

$$w_{n,p} = p^n + \sum_{j=1}^{n-1} \binom{n}{j} p^{n-j} q^j \max_{j \leq m \leq n-1} w_{m,p}, \quad w_{0,p} := 1. \quad (1)$$

Two strategies are natural: ONE sets aside a single head each round (value $a_{n,p}$), and ALL sets aside *every* head (value $b_{n,p}$). The companion paper [1] determines $w_{n,p}$ *exactly* at $p = \frac{1}{2}$, identifies the limit $W(p) = \lim_n w_{n,p}$ for $p > \frac{1}{2}$, and gives a first-order perturbation expansion for $p < \frac{1}{2}$ near $\frac{1}{2}$. All of those results are machine-checked in Lean 4.

The present paper concerns the regime left open by a perturbative analysis: the genuine large- n behaviour for $p < \frac{1}{2}$, which turns out to be *oscillatory* rather than convergent.

1.2 Main results

We state the two theorems at once; everything else in the paper is in their service.

Theorem 1.1 (oscillation of the optimal value). *For every $p \in (0, \frac{1}{2})$ the sequence $n \mapsto w_{n,p}$ is not eventually monotone. It has infinitely many strict local maxima and infinitely many strict local minima.*

Theorem 1.2 (log-periodicity of strategy ALL). *For every $p \in (0, \frac{1}{2})$,*

$$b_{n,p} = \Phi_p(\log n) (1 + o(1)) \quad (n \rightarrow \infty),$$

where Φ_p is continuous, strictly positive, non-constant, and periodic of period $L := \log(1/q)$. In particular $b_{n,p}$ does not converge: $\liminf_n b_{n,p} < \limsup_n b_{n,p}$.

Theorem 1.1 sharpens, and corrects the natural first guess about, the first-order picture of [1]: that analysis shows $w_{n,p}$ to be eventually increasing *at first order in $\frac{1}{2} - p$* , but the exact sequence is not eventually monotone at all. The oscillation is non-perturbative — its amplitude is exponentially small in $1/(\frac{1}{2} - p)$ — which is why it is invisible to any finite-order expansion.

1.3 Strategy of proof

The proof of Theorem 1.1 has a rigid skeleton, which we record now so that the role of each later section is clear. For a real sequence (x_n) put $d_n := x_{n+1} - x_n$.

Step 0 (reduction). A bounded sequence has infinitely many strict local maxima and minima as soon as it is *not eventually monotone*, i.e. as soon as $d_n > 0$ infinitely often and $d_n < 0$ infinitely often (Lemma 5.1). Thus Theorem 1.1 splits into

(A) $(w_{n,p})$ is not eventually non-decreasing, (B) $(w_{n,p})$ is not eventually non-increasing.

Step A (the increasing side). On a stretch where w is non-decreasing, the suffix-maximum in (1) collapses to its right endpoint and the Bellman equation becomes the linear strategy-ONE recursion, whose one-step increment has the sign-definite dominant term $-q^n w_{n-1}$. Since $q > p$, this forces an eventual decrease. Statement (A) follows (Proposition 4.1); the argument is elementary and complete.

Step B (the decreasing side). On a stretch where w is non-increasing, the suffix-maximum collapses to its *left* endpoint, and w becomes — exactly — a strategy-ALL sequence (Lemma 5.2). Strategy ALL oscillates (Theorem 1.2); hence w cannot have a non-increasing tail. The transfer of the oscillation from ALL to w is carried out by a renormalisation argument (Section 6): a windowed fundamental coefficient $A(t)$ is shown to satisfy a propagation identity across the geometric scale, from which non-convergence — hence non-monotonicity — follows.

The asymmetry between Steps A and B is structural: the increasing side reduces to a contractive linear recursion with a sign-definite drift, the decreasing side to a scale-invariant recursion whose solutions oscillate. Step A is elementary; Step B carries the analytic weight of the paper.

2 Setup and notation

Throughout, $p \in (0, \frac{1}{2})$ is fixed, $q = 1 - p \in (\frac{1}{2}, 1)$, and we use freely that $q > p$, $p/q < 1$ and $2p < 1$. We write $L := \log(1/q) > 0$ and $\omega := 2\pi/L$.

The value $w_{n,p}$ satisfies (1); it is a probability, so $0 < w_{n,p} \leq 1$, and in fact $w_{n,p} < 1$ for $n \geq 1$. The two reference strategies satisfy *linear* recursions:

$$a_{n,p} = p^n + (1 - p^n - q^n) a_{n-1,p}, \quad b_{n,p} = \sum_{j=0}^{n-1} \binom{n}{j} p^{n-j} q^j b_{j,p}, \quad (2)$$

with $a_{0,p} = b_{0,p} = 1$, and $w_{n,p} \geq \max(a_{n,p}, b_{n,p})$ since the optimum dominates any fixed strategy. We use the binomial identities $\sum_{j=1}^{n-1} \binom{n}{j} p^{n-j} q^j = 1 - p^n - q^n$ and $\sum_{j=0}^{n-1} \binom{n}{j} p^{n-j} q^j = 1 - q^n$. Where the value of p is fixed we abbreviate $w_n := w_{n,p}$, $b_n := b_{n,p}$.

3 Strategy All: log-periodic asymptotics

This section proves Theorem 1.2. We summarise the analysis; the full development is in the companion note `strategy_all_asymptotics.tex`.

3.1 The functional equation

Let $B(z) := \sum_{n \geq 0} b_n z^n / n!$ be the exponential generating function of strategy ALL. The convolution structure of (2) yields the first-order q -difference equation

$$B(z) = 1 + (e^{pz} - 1) B(qz), \quad (3)$$

and iterating it (the remainder tends to 0 since each factor $e^{pq^k z} - 1 \rightarrow 0$),

$$B(z) = \sum_{M \geq 0} \prod_{k=0}^{M-1} (e^{pq^k z} - 1). \quad (4)$$

Extracting coefficients exhibits b_n/p^n as a q -analogue of the Fubini (ordered-Bell) polynomials; equivalently, with $X_1, \dots, X_n$ i.i.d. geometric, $b_{n,p} = \Pr(\{X_1, \dots, X_n\} = \{1, \dots, \max_i X_i\})$.

3.2 Saddle-point and Mellin analysis

Writing $e^{pq^k z} - 1 = e^{pq^k z} (1 - e^{-pq^k z})$ and using $p \sum_{k=0}^{M-1} q^k = 1 - q^M$,

$$\prod_{k=0}^{M-1} (e^{pq^k z} - 1) = e^{z(1-q^M)} \prod_{k=0}^{M-1} (1 - e^{-pq^k z}).$$

On the positive axis the terms of (4) are dominated by M near $K(z) := \log(pz)/\log(1/q)$; a saddle-point evaluation gives $B(z) = e^z \Phi(\log z)(1 + o(1))$ with Φ bounded and periodic of period L , and a Hayman-type transfer to coefficients yields $b_{n,p} = \Phi_p(\log n)(1 + o(1))$. A Mellin analysis of the harmonic sum $\sum_k \log(1 - e^{-pq^k z})$ identifies the Fourier coefficients of Φ_p at frequencies $s_\ell = 2\pi i \ell / L$ as

$$c_\ell = -\frac{1}{L} p^{-s_\ell} \Gamma(s_\ell) \zeta(s_\ell + 1), \quad \ell \in \mathbb{Z}.$$

For $\ell \neq 0$ these are non-zero: Γ has no zeros and $\zeta(1 + it) \neq 0$ for real $t \neq 0$ (the non-vanishing of ζ on the line $\Re s = 1$). Hence Φ_p is non-constant and $b_{n,p}$ does not converge — which is Theorem 1.2.

► **OPEN GAP — Gap A — Hayman regularity. Work continues here.**

The Hayman-type transfer from $B(z) \sim e^z \Phi(\log z)$ to the coefficient asymptotics $b_n \sim \Phi_p(\log n)$ requires admissibility / sector control of $B(z)$. Here B is entire of order 1 and type 1 (indeed $B(r) \leq e^r$, since $0 \leq b_n \leq 1$); the delicate feature is the non-constant log-periodic factor, which the closure theory of H-admissible functions does not cover. A direct sector saddle-point estimate is needed.

► **OPEN GAP — Gap B — Mellin contour. Work continues here.**

The Mellin step — the strip of analyticity and the contour shift for the transform of $\log B(z) - z$ — is standard in spirit (Flajolet–Sedgewick) but is not carried out here.

Remark 3.1. Gaps A and B concern the *quantitative* asymptotics. The qualitative conclusion of Theorem 1.2 that we actually use downstream — that $b_{n,p}$ does not converge, equivalently Φ_p non-constant — is, granted the Mellin framework, secured by $\zeta(1+it) \neq 0$.

4 Step A: $w_{n,p}$ is not eventually increasing

Proposition 4.1. *For $p \in (0, \frac{1}{2})$ there is no M with $w_n \geq w_{n-1}$ for all $n > M$. That is, $w_n < w_{n-1}$ for infinitely many n .*

Proof. Suppose $w_n \geq w_{n-1}$ for all $n > M$, i.e. w is non-decreasing on $\{M, M+1, \dots\}$. Fix $n > M+1$ and split the Bellman sum (1) at $j = M$, writing $S_{j,n} := \max_{j \leq m \leq n-1} w_m$. For $j \geq M$ the index window lies in the non-decreasing range, so $S_{j,n} = w_{n-1}$; for every j one has $w_{n-1} \leq S_{j,n} \leq 1$. With $U_n := \sum_{j=1}^{M-1} \binom{n}{j} p^{n-j} q^j$ and $T_n := \sum_{j=1}^{M-1} \binom{n}{j} p^{n-j} q^j S_{j,n}$,

$$w_n = p^n + T_n + (1 - p^n - q^n - U_n) w_{n-1},$$

hence

$$w_n - w_{n-1} = p^n(1 - w_{n-1}) - q^n w_{n-1} + (T_n - U_n w_{n-1}).$$

Since $w_{n-1} \leq S_{j,n} \leq 1$ one has $0 \leq T_n - U_n w_{n-1} \leq U_n$, and $U_n \leq (M-1)n^{M-1}p^{n-M+1} = C_1 n^{M-1}p^n$ with $C_1 = (M-1)p^{1-M}$. With $c := w_M > 0$ and $w_{n-1} \geq c$,

$$w_n - w_{n-1} \leq (1 + C_1 n^{M-1})p^n - c q^n.$$

As $p/q < 1$, the ratio $(1 + C_1 n^{M-1})(p/q)^n \rightarrow 0$, so the right-hand side is negative for all large n — contradicting $w_n \geq w_{n-1}$. $\square$

Proposition 4.1 is statement (A) of Section 1.3 and is fully rigorous.

5 Step B, part 1: reduction and the transfer lemma

Lemma 5.1 (reduction). *A bounded real sequence (x_n) that is neither eventually non-decreasing nor eventually non-increasing has infinitely many strict local maxima and infinitely many strict local minima, provided $x_n \neq x_{n+1}$ for all large n .*

Proof. Not eventually non-decreasing means $d_n := x_{n+1} - x_n < 0$ infinitely often; not eventually non-increasing means $d_n > 0$ infinitely often. Between a sign change of d from $-$ to $+$ lies a local minimum, between $+$ and $-$ a local maximum; infinitely many sign changes yield infinitely many of each, strict under $x_n \neq x_{n+1}$. $\square$

Lemma 5.2 (transfer to strategy ALL). *Suppose $w_n \geq w_{n+1}$ for all $n \geq M$. Define $x_0 := 1$, $x_j := \max_{j \leq m \leq M} w_m$ for $1 \leq j \leq M-1$, and $x_j := w_j$ for $j \geq M$. Then $x_n = w_n$ for $n \geq M$, and (x_n) satisfies the exact strategy-ALL recursion*

$$x_n = \sum_{j=0}^{n-1} \binom{n}{j} p^{n-j} q^j x_j \quad \text{for every } n > M.$$

Proof. Fix $n > M$. In the Bellman equation (1) the suffix-maximum $S_{j,n} = \max_{j \leq m \leq n-1} w_m$ collapses to the *left* endpoint on the non-increasing range: for $M \leq j \leq n-1$, $S_{j,n} = w_j = x_j$; for $1 \leq j \leq M-1$, $S_{j,n} = \max(\max_{j \leq m \leq M-1} w_m, w_M) = \max_{j \leq m \leq M} w_m = x_j$, independent of n . Writing the p^n term as $\binom{n}{0} p^n q^0 x_0$ gives the displayed recursion. $\square$

Lemma 5.2 is the crux of Step B: it says that a non-increasing tail of w is, exactly, a strategy-ALL sequence with finitely many initial values modified. The remaining task is to show such a sequence cannot exist, by transferring the oscillation of Theorem 1.2.

6 Step B, part 2: the renormalisation argument

Throughout this section (x_n) is a bounded solution of the strategy-ALL recursion for $n > M$ — either b itself, or a transferred sequence of Lemma 5.2.

6.1 The recursion as a smoothing in log-scale

With $J_n \sim \text{Bin}(n, q)$ the recursion reads $x_n = E[x_{J_n}]/(1 + q^n)$. Passing to $t := \log n$ and the log-linear interpolation y of $n \mapsto x_n$, and using $\log(qn) = \log n - L$, one application of the recursion shifts the log-scale down by exactly one period L and smooths by a kernel of width $\sigma(t) \sim \sqrt{p/q} e^{-t/2} \rightarrow 0$.

6.2 The local fundamental coefficient

Fix a window $\chi \in C_c^\infty(\mathbb{R})$ with $\int \chi = 1$ and $\hat{\chi}(2\pi) = 0$, and set

$$A(t) := \frac{1}{L} \int_{\mathbb{R}} y(s) e^{-i\omega s} \chi\left(\frac{s-t}{L}\right) ds. \quad (5)$$

The condition $\hat{\chi}(2\pi) = 0$ makes A annihilate the constant (mean-level) component, so A extracts the fundamental oscillation. If x_n converges then $A(t) \rightarrow 0$.

6.3 The propagation estimate

Proposition 6.1 (propagation). *There are $\kappa(t) \in \mathbb{C}$ and a remainder $r(t)$ with*

$$A(t+L) = \kappa(t) A(t) + r(t), \quad |\kappa(t) - 1| = O(e^{-t}), \quad |r(t)| = O(e^{-t/2}).$$

The factor κ is fully under control.

Lemma 6.2 (the factor κ). *For $J_n \sim \text{Bin}(n, q)$, $\kappa(n) = E[(J_n/(qn))^i \mathbf{1}_{J_n \geq 1}] = 1 - \frac{p}{2qn}(\omega^2 + i\omega) + O(n^{-3/2})$. In particular $|\kappa - 1| = O(1/n) = O(e^{-t})$.*

Proof (sketch). Put $V_n = J_n/(qn) - 1$; the binomial central moments give $E[V_n^2] = p/(qn)$ and $E[|V_n|^3] = O(n^{-3/2})$. Expanding $(1 + V_n)^{i\omega}$ to second order on $\{|V_n| \leq \frac{1}{2}\}$ and bounding the tail by Hoeffding yields the claim. (Full proof: the development note, Lemma 8.4.) $\square$

For the remainder we record the structural identities on which the estimate rests. Both are exact.

Lemma 6.3 (windowed-smoothing identities). *Let $A[y]$ be given by (5).*

- (a) $A[y] \in C^\infty$ in t , with $\partial_t A[y] = -\frac{1}{L} A^{[X']}[y]$; since $\widehat{\chi}'(2\pi) = 2\pi i \widehat{\chi}(2\pi) = 0$, every t -derivative of $A[y]$ is again a constant-annihilating windowed coefficient, hence bounded.
- (b) For a probability law ν and $S_\nu y(s) := \int y(s + \zeta) \nu(d\zeta)$,

$$A[S_\nu y](t) = \int e^{i\omega\zeta} A[y](t + \zeta) \nu(d\zeta).$$

Proof. (a) Differentiate under the integral. (b) Fubini, then substitute $s \mapsto s + \zeta$ in the inner integral. $\square$

Decomposition of the remainder. Applying the recursion inside (5) at $t+L$ and substituting $s \mapsto s+L$ (using $e^{-i\omega L} = 1$),

$$A(t+L) = \frac{1}{L} \int E_{\nu^{(s+L)}}[y(s+\zeta)] e^{-i\omega s} \chi\left(\frac{s-t}{L}\right) ds + \varepsilon_{\text{tail}}(t),$$

where $\nu^{(s+L)}$ is the law of $\log(J_{e^{s+L}}/(qe^{s+L}))$. Four contributions arise:

- (i) $\varepsilon_{\text{tail}}$ (the factor $1/(1+q^n)$ and the indices $j \leq M$): super-exponentially small. *Rigorous.*
- (ii) the discretisation error (sum versus integral): with a C^∞ compactly supported window, super-polynomially small by Poisson summation. *Rigorous.*
- (iii) the main term: replacing $\nu^{(s+L)}$ by the window-centre law $\nu^{(t+L)}$ and applying Lemma 6.3(b) and a Taylor expansion of $A(t+\zeta)$ (legitimate by Lemma 6.3(a)), together with Lemma 6.2, gives $\kappa(t)A(t) + O(e^{-t/2})$. *Rigorous.*
- (iv) the residual $\rho(t) = \frac{1}{L} \int e^{-i\omega s} \chi\left(\frac{s-t}{L}\right) \int y(s+\zeta) (\nu^{(s+L)} - \nu^{(t+L)})(d\zeta) ds$, measuring the slow variation of the smoothing kernel across the window.

Thus $r(t) = \varepsilon_{\text{tail}} + \varepsilon_{\text{disc}} + O(e^{-t/2}) + \rho(t)$, and Proposition 6.1 follows once $|\rho(t)| = O(e^{-t/2})$.

► OPEN GAP — Gap C — the remainder ρ . Work continues here.

The bound $|\rho(t)| = O(e^{-t/2})$ is the one remaining piece of Proposition 6.1; the other three contributions are established above. A naive total-variation bound fails ($\|\nu^{(s+L)} - \nu^{(t+L)}\|_{\text{TV}} = \Theta(1)$ across the window); the cancellation is genuine and is supplied by the windowed projection onto the fundamental mode. Closing it requires a density-level local central limit theorem for $\log J_n$. This is where work continues.

6.4 From propagation to non-convergence

Theorem 6.4 (renormalisation). *Let (x_n) be a bounded solution of the strategy-ALL recursion for $n > M$. Then for each t_0 the sequence $A(t_0 + kL)$ converges, as $k \rightarrow \infty$, to a limit A_∞ . Moreover: if x_n converges then $A_\infty = 0$; and if $A_\infty \neq 0$ then (x_n) is not eventually monotone.*

Proof. Write $A_k = A(t_0 + kL)$, $\kappa_k = \kappa(t_0 + kL)$, $r_k = r(t_0 + kL)$. By Proposition 6.1, $\sum_k |\kappa_k - 1| < \infty$ and $\sum_k |r_k| < \infty$ (geometric series in $e^{-L} < 1$). Solving $A_{k+1} = \kappa_k A_k + r_k$,

$$A_k = \left(\prod_{j < k} \kappa_j \right) A_0 + \sum_{j < k} \left(\prod_{j < i < k} \kappa_i \right) r_j \rightarrow A_\infty,$$

the products and series converging absolutely. If x_n converges then $A(t) \rightarrow 0$, so $A_\infty = 0$. If $A_\infty \neq 0$ then $A_k \not\rightarrow 0$, so x_n does not converge; a bounded non-convergent sequence is not eventually monotone. $\square$

7 The base case and the proof of Theorem 1.1

Theorem 6.4 reduces Step B to a single non-degeneracy: the limit A_∞ must be non-zero.

Proposition 7.1 (base case). *For every transferred sequence (x_n) of Lemma 5.2, $A_\infty \neq 0$.*

By linearity $A_\infty = A_\infty(b) + A_\infty(g)$, where $g = x - b$ is a non-negative solution of the homogeneous recursion (companion note, §7). Here $A_\infty(b)$ is the fundamental Fourier coefficient $\gamma_1(p)$ of the profile Φ_p of Theorem 1.2, which is non-zero by $\zeta(1 + it) \neq 0$; the modulation by g does not cancel it.

► **OPEN GAP — Gap D — the non-degeneracy. Work continues here.**

$A_\infty \neq 0$ is the genuine arithmetic heart of the proof. For the pure strategy ALL it is $\gamma_1(p) \neq 0$, which holds within the Mellin framework (Gaps A, B) via the non-vanishing of ζ on the line $\Re s = 1$. For a transferred sequence one must further exclude that the perturbation g cancels the oscillation; the companion note reduces this to the linear independence, modulo constants, of the profiles of a family of auxiliary functions B_r . Work continues here.

Proof of Theorem 1.1. Fix $p \in (0, \frac{1}{2})$. By Lemma 5.1 it suffices to show that $(w_{n,p})$ is neither eventually non-decreasing nor eventually non-increasing.

Not eventually non-decreasing is Proposition 4.1.

Not eventually non-increasing. Suppose, for contradiction, $w_n \geq w_{n+1}$ for all $n \geq M$. By Lemma 5.2 the transferred sequence (x_n) satisfies $x_n = w_n$ for $n \geq M$ and solves the strategy-ALL recursion for $n > M$; being a tail of a monotone bounded sequence, x_n converges. By Theorem 6.4 then $A_\infty = 0$, contradicting Proposition 7.1. Hence no non-increasing tail exists.

Both one-sided statements hold, so $(w_{n,p})$ is not eventually monotone; Lemma 5.1 gives infinitely many strict local maxima and minima. $\square$

Remark 7.2 (strictness). Lemma 5.1 delivers *strict* extrema once $w_{n,p} \neq w_{n+1,p}$ for all large n , which holds in all computed data and is expected for every $p \in (0, \frac{1}{2})$; absent it, Theorem 1.1 holds with the non-strict notion of local extremum.

8 Numerical evidence

The two predictions have been checked numerically in arbitrary-precision arithmetic (`simulation/b_asymptotics.py`, `simulation/renorm_validation.py`). For strategy ALL the log-periodic oscillation of Theorem 1.2 is confirmed across fourteen orders of magnitude in its amplitude. For the renormalisation argument: the windowed coefficient $A(t_0 + kL)$ is observed to converge to a non-zero limit A_∞ for every tested $p \in \{0.35, 0.40, 0.42, 0.45\}$, and $n |\kappa(n) - 1|$ converges to the constant $\frac{p}{2q} |\omega^2 + i\omega|$ of Lemma 6.2 — a direct quantitative confirmation of Proposition 6.1.

9 Status and open problems

The skeleton of the proof is complete and the rigorous backbone is substantial: the reduction (Lemma 5.1), the elementary Step A (Proposition 4.1), the exact transfer lemma (Lemma 5.2), the windowed-smoothing identities (Lemma 6.3), the factor estimate (Lemma 6.2), three of the four contributions to the propagation remainder, and the renormalisation theorem (Theorem 6.4). Four steps remain open; they are flagged in the text and summarised here.

Gap	What is missing	Where
A	Hayman regularity / sector control of $B(z)$	§3
B	Mellin contour justification	§3
C	the remainder bound $ \rho(t) = O(e^{-t/2})$	§6
D	the non-degeneracy $A_\infty \neq 0$	§7

A complete proof of Theorem 1.1 consists of Proposition 4.1 (done) together with Gaps C and D; Gaps A and B underpin Gap D and the quantitative form of Theorem 1.2. The natural order of attack is C, then A and B, then D. Gap C has been reduced (§6) to a single density-level local limit theorem for the binomial; that is the current front of work.

References

- [1] P. Pfaffelhuber, *Optimal strategies in the all-heads coin game*, companion paper.
- [2] E. A. van Doorn, *Optimal play of a coin-flipping game*, preprint, arXiv:2406.14700, 2024.
- [3] P. Flajolet, R. Sedgewick, *Analytic Combinatorics*, Cambridge University Press, 2009.